

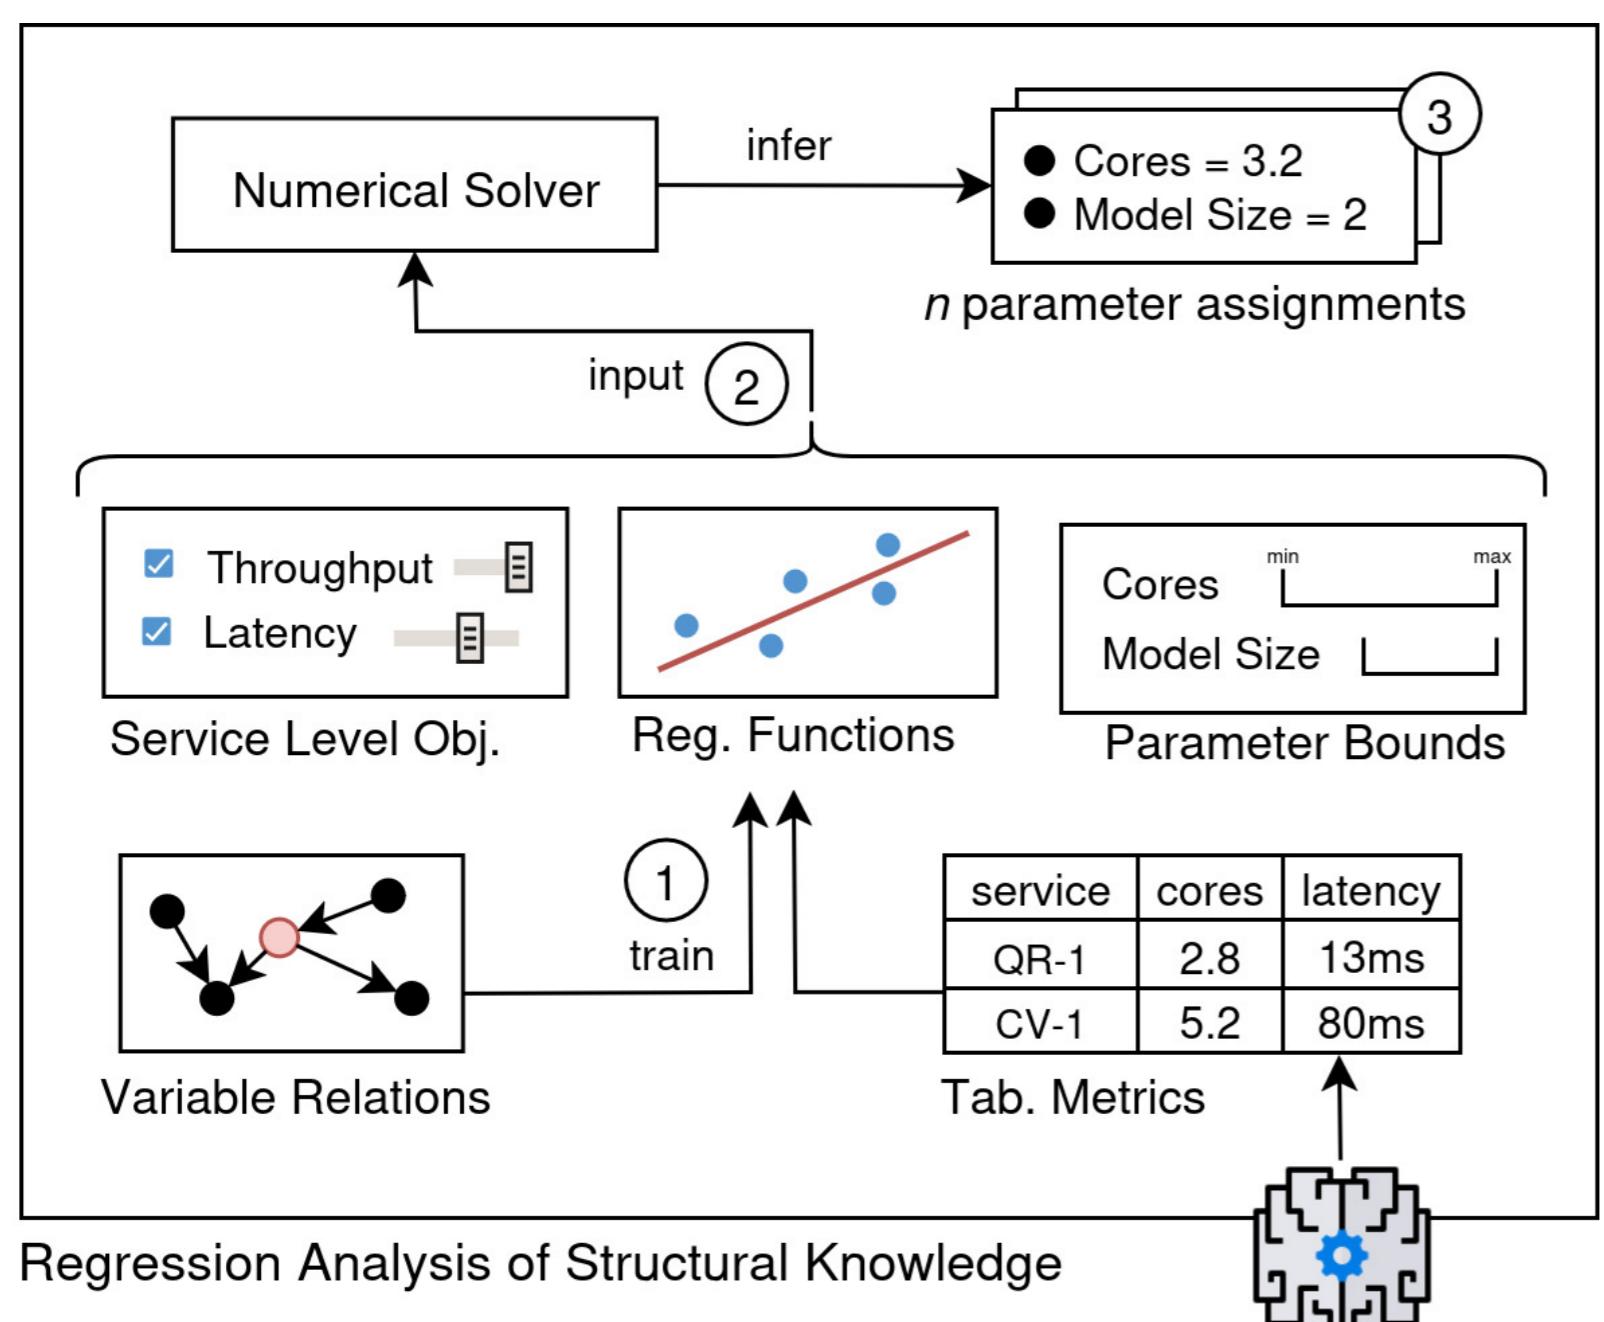
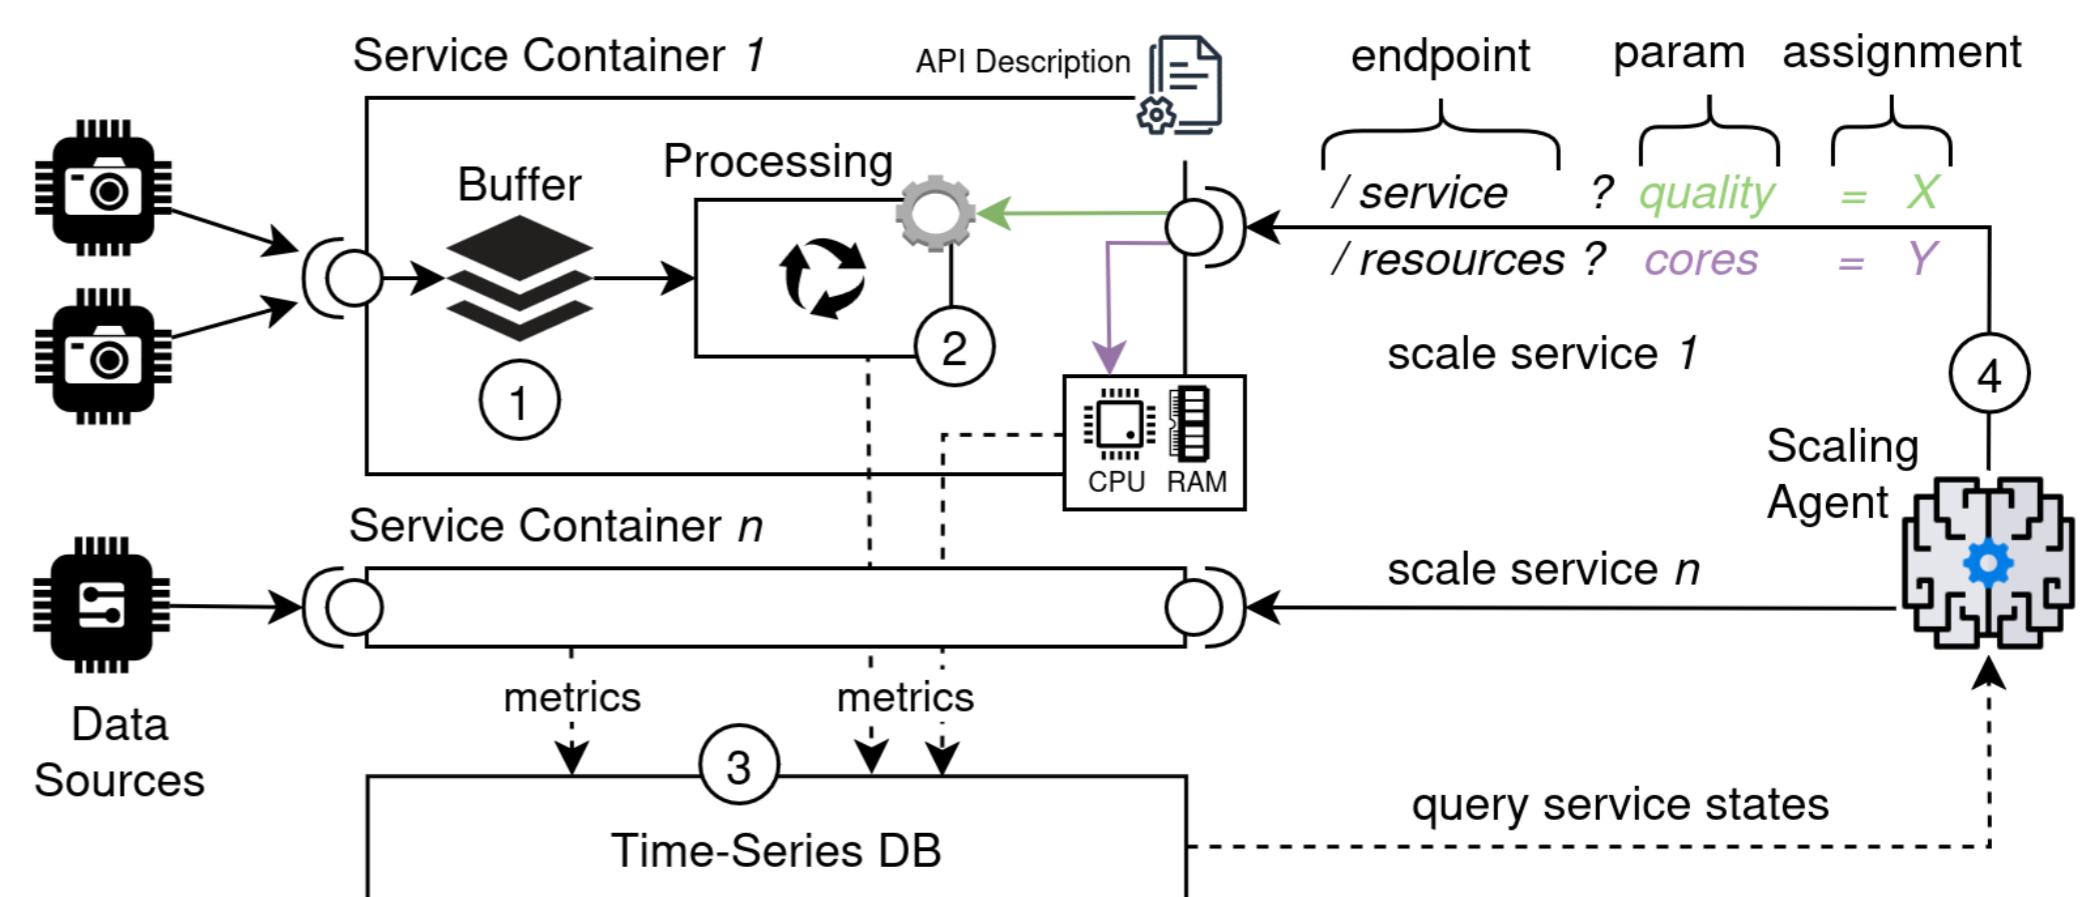
How to Gradually Improve Autoscaling through Analysis of Structural Knowledge

Autoscaling Environment

Containerized service execution on **one** host; scale services through **multiple** service/device params.

Service execution: (1) buffer and (2) process data, and (3) log service and device metrics in a DB.

Scaling agent: (4) develop a scaling policy and acts on service and device through REST API.



Scaling Agent Cycle

Analyzes **structural knowledge** to infer optimal scaling policy under custom application SLOs and parameter bounds.

- 1) Using metrics and **variable relations**, estimate state transitions through regression functions. Quantifies impact of **interventionable parameters** on global SLO fulfillment.
- 2) Supply regression functions, SLOs, and parameter bounds to numerical solver; through **gradient descent** find optimal assignment for multiple services under resource limitations.
- 3) Use exposed REST API to scale the services parameters and their containerized resources accordingly.

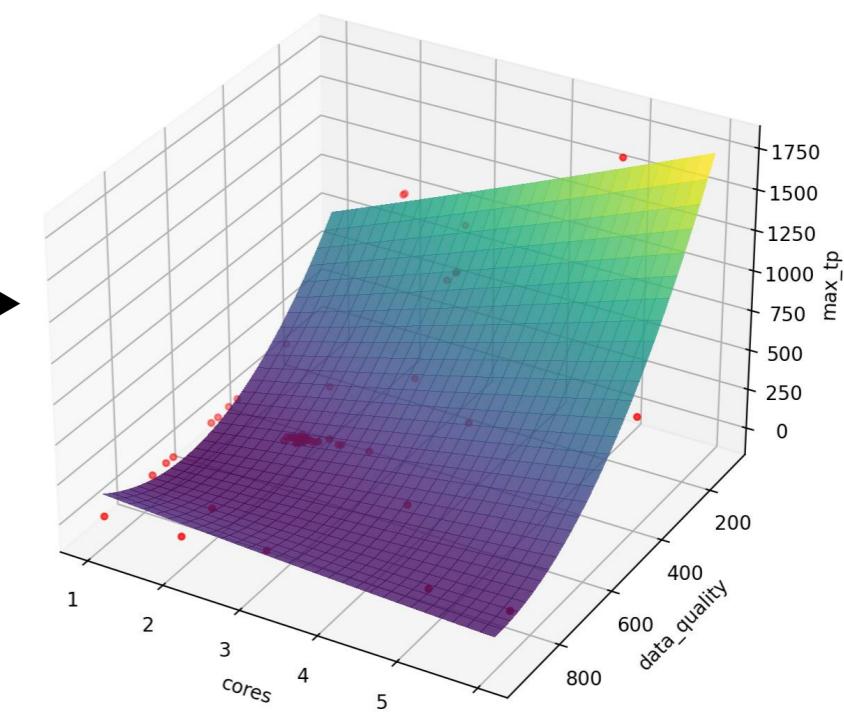
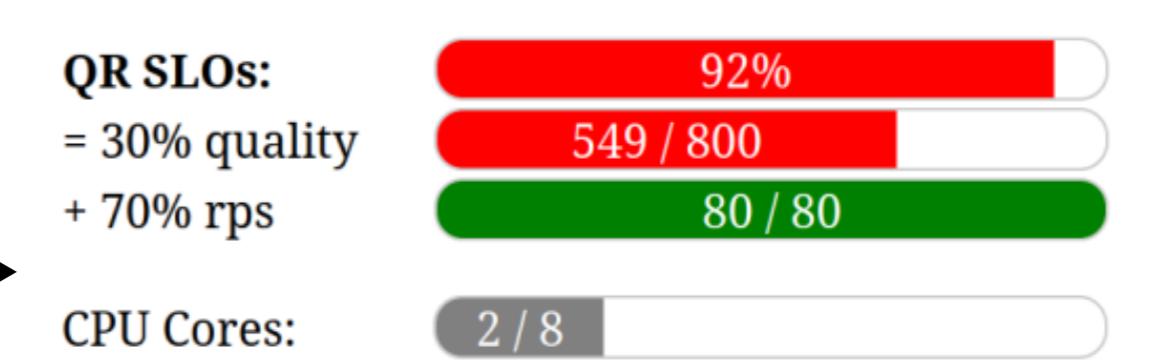
Demo Contents

Scaling agent needs to **optimize** the global SLO fulfillment of three processing services that compete for 8 CPU cores. Agent operates in cycles of 10s, continuously executing above cycle. For 30 iterations, agent **explores**, afterwards exploits.

Service output: Current output of three processing services, i.e., detected QR codes, Yolov8 objects, and Lidar mobile map

Service states and SLOs: For interventionable and dependent variables, show SLO fulfillment relative to thresholds; overall SLO fulfillment **weighted** according to preferences. CPU cores are allocated between services within global limitations.

Regression model: Show impact of interventionable parameters on dependent variable (i.e., **expected rps**). Improve over time.



**Demo
Video**



**Project
Repo**